

CS 342 Design Document

Project 2D: Implementing System Calls

Rajat Kateja (r.kateja@iitg.ernet.in)
Mathematics and Computing
10012338

TA: Lalatendu Behera
Chetti Prasad

----DATA STRUCTURES----

The following data structures were introduced:

In thread.h

Added the following to struct thread

```
struct list fd_list;           /* File Descriptor List */
struct lock child_lock;       /* Lock for the parent to wait on child */
bool child_load_failed;       /* Boolean to store if the child load failed */
struct list dead_child_list;   /* Store a list of dead children of a process */
bool waited_on_by_parent;     /* Did parent thread already wait on this thread */
```

Introduced the following data structure:

Structure fd is the file descriptor which has three fields, first is an integer related to each open file for each thread, next is the file pointer, and last is the list elem so that the instances of this structure can be inserted into the struct list fd_list of each thread.

```
struct fd
{
    int file_no;
    struct file *file;
    struct list_elem elem;
};
```

Structure dead_thread stores the dead child thread of a parent thread which may be used by the parent thread to wait on the child thread (i.e. return with the child's exit status) even after the child thread has exited. It has three fields, first is the tid of the thread to identify the thread, next is the exit status of the thread and last is the list_elem so that instances of this struct can be inserted into the dead_child_list of each of the threads.

```
struct dead_thread
{
    int tid;
    int exit_status;
    struct list_elem elem;
};
```

----ALGORITHMS----

The following algorithms were used to implement the various system calls:

First of all in the syscall handler, the `f->esp` pointer is checked to be non NULL, must belong to the user valid address space and must be a valid page in the page table of the user program (checked by `lookup_page` function). If any of these tests fails, then the thread is exited with a -1 exit status.

Wrote a function **`create(int no_of_args, struct frame *f)`**. This function checks that the `no_of_args` arguments in the frame stack are all valid, the three tests mentioned above are applied to each of them, and if any of that fails, it returns false. In each of the system calls, depending on the number of arguments that the system call receives, the function is called along with the frame. If the return is false, **`quit`** function is called (implemented for this project) which makes the thread exit with exit status -1.

Halt

This system call simply switches off the system by calling `power_off()` function.

Exit

It takes the exit status of the thread, prints that the thread is exiting with the given exit status in the required format, sets the return value to the exit status, and then it pushes the thread into its parents `dead_thread_list`, so that it can be waited upon by the parent even after it is dead. Finally `thred_exit()` is called to make the thread exit.

Exec

This function is used to create new child functions, the basic functionality achieved by calling the `process_execute()` function with the received argument of the process, which does the required job of creating the new thread and putting it in the ready queue. Then the function puts the parent thread in the `child_lock` waiting list, and blocks itself. This is used as a synchronisation because it is required that the parent cannot return from this function until the child has successfully loaded. So now the parent goes into block state and is in the `child_lock` lock's waiting list. From here, the parent is woken up by the child only after its loading is finished (in `process.c` `start_process` function). Here only, the requirement that is child loading has failed, the function must return -1 is met. If the loading of the child has failed, in the `start_process` function the child thread sets the `child_load_failed` Boolean field in the parent thread to be true. Next when the thread is unblocked from the `child_lock` waiting list, it checks for this field. If it is set to be true, the return value is set to be -1. Else the return value is the value returned by `process_execute` function which is nothing but the exit status of the child thread.

Wait

This system call is used to wait on a child thread by the parent thread. This functionality is basically achieved by calling the `process_wait()` function in `process.c` where all the requirements are taken care of. The following is performed in the `process_wait` function:

- First the `tid` is checked, in the `all_list` (where all threads are stored) by using the function `get_thread_by_id` which was made for the previous implementation as well (project 2B and 2C). If a valid thread is not found in the list, then the calling thread (i.e. parent thread's) `dead_child_list` is checked. If the thread is found there, the return value is set to the exit status of the thread, the child thread is removed from the list (since it has already been

waited upon once), the the function returns. If the thread is not found even in the `dead_thread_list`, then the function returns with -1 as required.

- If the thread was found in the `all_list` of thread, then it is checked that the `tid` must be of one of the parent's child threads (by iteratively checking the parents `child_list`). If thread is not found in the `child_list`, function returns with -1. Also if the thread has already been waited upon by the parent (checked by using the bool `waited_on_by_parent` field of the thread), then also the function returns with -1 status. All theses are as per the requirements specified in the pintos manual.
- If till this stage, the function has not returned, means we need to now wait on the given child thread. First its `waited_on_by_parent` is set to true (for future reference), and the parent's `waiting_on_child` is set to be this child thread. The the parent thread is blocked.
- The parent thread is then woken up by the child thread in its `thread_exit` function, by checking that the parent is non-NULL, and its `waitin_on_child` field.
- Once the thread wakes up, it checks its `dead_child_list` and removes this thread from the list. This step is done because before waking up the parent the child thread puts itself in the parent's `dead_child_list` (in the exit function), and this thread has already been waited upon, so it must not be in the `dead_child_list`.
- Finally the parent thread returns from the function.

Create

This system call is used to create a new file in the file system. This is done using the function `filesys_create()` provided in the `filesys.c` file using the appropriate arguments as passed to this function. The return value is also the same as the one returned by the mentioned function.

Remove

This function is the opposite of create function, used to remove a file from the file system. This is achieved using the `filesys_remove` function with the appropriate arguments.

Open

This system call opens the file specified in the file name received as the argument. It first performs the routine checks on the arguments, and returns with exit status -1, if any of the tests fails. Then it opens the file using `filesys_open` function which returns a file pointer. Next task is to put this file in the file descriptor of the thread. For this, first the `fd_list` of the thread is checked. If it is empty, means this is the first file that is being opened by the thread. In this case, the `file_no` of this file descriptor is set to 3 (since 0,1 and 2 are reserved). If the list is non-empty, then the `file_no` is set to be one larger than the previous `file_no`. Finally it is checked that the file opened is not the thread's executable or its parent's executable. If either of the above is true, then the file is rendered non-writeable using the `file_deny_write()` function.

Filesize

This function receives a file number as argument. This file number is transformed to file pointer from the `fd_list` of the thread using the `get_fd_by_no()` function (implemented). Then the `file_length()` function provided in `file.c` is used to get the size of the file, which is returned.

Read

This function also received a file number as argument (along with size and buffer). This number is used to get the file descriptor from the `fd_list` of the thread using the `get_fd_by_no()` function. Then if the number is 0, which is the case of `stdin`, then `input_getc()` function is used. In other cases `file_read` function is used by passing the appropriate arguments.

Write

This function is similar to the read function, only change is that it is used to write on the file provided by the number. Like before first the file descriptor is found, then the appropriate function is called, i.e. putbuf in case of stdout and file_write in all other cases.

Seek

This is used to set the read/write position in the file. It is implemented using the file_seek function provided in file.c.

Tell

It is used to get the read/write position in the file, which is found from the file_tell function.

Close

This is used to close the file preciouslt opened. Firstly the file pointer is found from the file descriptor of the file (found using get_fd_by_no function), then the instance of this file is removed from the fd_list of the thread and then the file is closed using the file_close function.

One particular implementation is in the **exception.c** file, where in case of a page fault exception, if the page faule was cause due to a NULL pointer being dereferenced or the pointer not being a valid used addredd, then the thread quits with -1 status by calling the quit function. This is done to handle the cases where a bad read or write is attempted.

---SYNCHRONISATION---

A parent must not return from the exec system call until the child has been loaded (irrespective whether successfully or unsuccessfully). This is achieved by using the child_lock. The parent is inserted in this lock's waiting list and blocked. Then the child acquires the lock in the start of it's start_process and releases the lock only after it has returned from the load function, thereby unblocking the parent thread.

---RATIONALE---

1. The code is perfectly synchronized,so even in case of arbitrary switches the deadlock situation will not occur.
2. Since it does not pass the multi-oom test,it is likely that the in case of overload of memory kernel may panic or crash (since the process may run out of memory) .

---CODE SNIPPETS-----

Here I am providing the code snippets of the system calls, and the functions written in the syscall.c file.

```
static void
syscall_handler (struct intr_frame *f)
{
    /*****
    Implementing system calls for project 2B
    *****/
}
```

```

struct thread *cur = thread_current();
uint32_t *valid = lookup_page(cur->pagedir, f->esp, false);
if(f->esp == NULL || !is_user_vaddr (f->esp) || valid == NULL)
    quit(f);

int syscall = * (int *) (f->esp);

switch(syscall)
{
    case(SYS_HALT):
    {
        power_off();
        break;
    }
    case(SYS_EXIT):
    {
        if(!check(1, f))
            quit(f);
        int status = * (int *) (f->esp + 4);
        f->eax = status;
        printf("%s: exit(%d)\n", cur->name, status);
        struct dead_thread *thread_to_die = (struct dead_thread
*)malloc(sizeof(struct dead_thread));
        thread_to_die->exit_status = status;
        thread_to_die->tid = thread_current()->tid;
        list_push_front(&thread_current()->parent->dead_child_list,
&thread_to_die->elem);
        thread_exit();
        break;
    }
    case(SYS_EXEC):
    {
        if(!check(1, f))
            quit(f);
        char *process = * (char **) (f->esp + 4);
        if(process == NULL || !is_user_vaddr (process) || lookup_page(cur-
>pagedir, process, false) == NULL)
            quit(f);
        f->eax = process_execute(process);
        list_push_front(&thread_current()->child_lock.waiters, &thread_current()-
>elem);
        thread_block();
        if(thread_current()->child_load_failed == true)
            f->eax = -1;
        break;
    }
    case(SYS_WAIT):
    {
        if(!check(1, f))
            quit(f);
        int tid = * (int *) (f->esp + 4);

```

```

        f->eax = process_wait(tid);
        break;
    }
    case(SYS_CREATE):
    {
        if(!check(2, f))
            quit(f);
        char *file = * (char **) (f->esp + 4);
        if(file == NULL || !is_user_vaddr (file) || lookup_page(cur->pagedir, file,
false) == NULL)

            quit(f);
        int file_size = * (int *) (f->esp + 8);
        f->eax = filesys_create(file, file_size);
        break;
    }
    case(SYS_REMOVE):
    {
        if(!check(1, f))
            quit(f);
        char *file = * (char **) (f->esp + 4);
        if(file == NULL || !is_user_vaddr (file) || lookup_page(cur->pagedir, file,
false) == NULL)

            quit(f);
        f->eax = filesys_remove(file);
        break;
    }
    case(SYS_OPEN):
    {
        if(!check(1, f))
            quit(f);
        char *file_name = * (char **) (f->esp + 4);
        struct file *file;
        if(file_name == NULL || !is_user_vaddr (file_name) || lookup_page(cur-
>pagedir, file_name, false) == NULL)
            quit(f);
        file = filesys_open(file_name);
        if(file == NULL)
            f->eax = -1;
        else
        {
            if(!list_empty(&cur->fd_list))
            {
                int prev_fd_no = list_entry(list_back(&cur->fd_list), struct
fd, elem)->file_no;

                struct fd *new_fd = (struct fd *)malloc(sizeof(struct fd));
                new_fd->file_no = prev_fd_no+1;
                new_fd->file = file;
                list_push_back(&cur->fd_list, &new_fd->elem);
                f->eax = new_fd->file_no;
            }
            else

```

```

        {
            struct fd *new_fd = (struct fd *)malloc(sizeof(struct fd));
            new_fd->file_no = 3;
            new_fd->file = file;
            list_push_back(&cur->fd_list, &new_fd->elem);
            f->eax = new_fd->file_no;
        }
    }
    if(strcmp(file_name, thread_current()->name) == 0 || strcmp(file_name,
thread_current()->parent->name) == 0)
    {
        file_deny_write(file);
    }
    break;
}
case(SYS_FILESIZE):
{
    if(!check(1, f))
        quit(f);
    int fd_no = * (int *) (f->esp + 4);
    struct fd *cur_fd = get_fd_by_no(fd_no);
    if(cur_fd != NULL)
        f->eax = file_length(cur_fd->file);
    break;
}
case(SYS_READ):
{
    if(!check(3, f))
        quit(f);
    int fd_no = * (int *) (f->esp + 4);
    char *buffer = * (char **) (f->esp + 8);
    if(buffer == NULL || !is_user_vaddr (buffer) || lookup_page(cur->pagedir,
buffer, false) == NULL)
        quit(f);
    int size = * (int *) (f->esp + 12);
    struct fd *cur_fd = get_fd_by_no(fd_no);
    if(cur_fd != NULL)
    {
        f->eax = file_read(cur_fd->file, buffer, size);
    }
    else if(fd_no == 0)
    {
        input_getc();
        f->eax = size;
    }
    break;
}
case(SYS_WRITE):
{
    if(!check(3, f))
        quit(f);

```

```

        int fd_no = * (int *) (f->esp + 4);
        char *buffer = * (char **) (f->esp + 8);
        if(buffer == NULL || !is_user_vaddr (buffer) || lookup_page(cur->pagedir,
buffer, false) == NULL)
            quit(f);
        int size = * (int *) (f->esp + 12);
        struct fd *cur_fd = get_fd_by_no(fd_no);
        if(cur_fd != NULL)
        {
            f->eax = file_write(cur_fd->file, buffer, size);
        }
        else if(fd_no == 1)
        {
            putbuf(buffer, size);
            f->eax = size;
        }

        break;
    }
    case(SYS_SEEK):
    {
        if(!check(2, f))
            quit(f);
        int fd_no = * (int *) (f->esp + 4);
        unsigned position = * (unsigned *) (f->esp + 8);
        struct fd *cur_fd = get_fd_by_no(fd_no);
        if(cur_fd != NULL)
            file_seek(cur_fd->file, position);
        break;
    }
    case(SYS_TELL):
    {
        if(!check(1, f))
            quit(f);
        int fd_no = * (int *) (f->esp + 4);
        struct fd *cur_fd = get_fd_by_no(fd_no);
        if(cur_fd != NULL)
            f->eax = file_tell(cur_fd->file);
        break;
    }
    case(SYS_CLOSE):
    {
        if(!check(1, f))
            quit(f);
        int fd_no = * (int *) (f->esp + 4);
        struct fd *cur_fd = get_fd_by_no(fd_no);
        if(cur_fd != NULL)
        {
            list_remove(&cur_fd->elem);
            file_close(cur_fd->file);
        }
    }

```



```

                break;
            }
        }

        //-----added code ends-----//
    }
    struct fd * get_fd_by_no(int no)
    {
        struct thread *cur = thread_current();
        struct list_elem *e;
        struct fd *cur_fd;
        for(e=list_begin(&cur->fd_list);e!=list_end(&cur->fd_list);e=list_next(e))
        {
            cur_fd = list_entry(e, struct fd, elem);
            if(cur_fd->file_no == no)
                break;
        }
        if(e==list_end(&cur->fd_list))
            return NULL;
        else
            return cur_fd;
    }
    bool check(int no_of_args, struct intr_frame *f)
    {
        int i;
        for(i=0;i<no_of_args;i++)
        {
            if(!is_user_vaddr (f->esp+4*(i+1)) || lookup_page(thread_current()->pagedir, f->esp+4*(i+1), false) == NULL)
                return false;
        }
        return true;
    }
    void quit(struct intr_frame *f)
    {
        f->eax = -1;
        printf("%s: exit(%d)\n", thread_current()->name, -1);
        thread_exit();
    }

```